

## Warehouse Scale Computing

### 1. Amdahl's Law

1) You are going to train the image classifier with 50,000 images on a WSC having more than 50,000 servers. You notice that 99% of the execution can be parallelized. What is the speedup?

$$100/(99*(1/50000) + 1) = \sim 100x$$

### 2. Failure in a WSC

1) In this example, a WSC has 55,000 servers, and each server has four disks whose annual failure rate is 4%. How many disks will fail per hour?

$$55000 * 4 * 0.04 / 365 / 24 = \sim 1$$

2) What is the availability of the system if it does not tolerate the failure? Assume that the time to repair a disk is 30 minutes.  $1/(1 + 0.5) = .66$

### 3. Performance of a WSC

	Local	Rack	Array
DRAM latency (us)	0.1	100	300
Global hit rate	90%	9%	1%
DRAM bandwidth (MiB/sec)	20,000	100	10
Disk bandwidth (MiB/sec)	200	100	10

1) Calculate the AMAT of this WSC. What is vital for WSC performance?

locality of data to disk drives speed

$$.1(10^{-6}) + .1(100*(10^{-6}) + .1(300(10^{-6}))) = 13.1(10^{-6})s$$

$$\text{hmmmm the answer says } (.9*.1 + .09*100 + .01*300)(10^{-6}) = 12.09(10^{-6})...$$

2) How long does it take to transfer 1,000 MiB a) between disks within the server, and b) between DRAM within the rack? What can you conclude from this example?

a) 5

b) 10

Keeping thing local is important to keeping things fast.

### 4. Power Usage Effectiveness (PUE) = (Total Building Power) / (IT Equipment Power)

Sources speculate Google has over 1 million servers. Assume each of the 1 million servers draw an average of 200W, the PUE is 1.5, and that Google pays an average of 6 cents per kilowatt-hour for datacenter electricity.

1) Estimate Google's annual power bill for its datacenters.

$$\sim 157 \text{ M\$}$$

2) Google reduced the PUE of a 50,000 machine datacenter from 1.5 to 1.25 without decreasing the power supplied to the servers. What's the cost savings per year?

$$\sim 1.3 \text{ M\$}$$

## Map Reduce

Use pseudocode to write MapReduce functions necessary to solve the problems below. Also, make sure to fill out the correct data types. Some tips:

- The input to each MapReduce job is given by the signature of the **map()** function.
- The function **emit(key k, value v)** outputs the key-value pair **(k, v)**.
- The **for(var in list)** syntax can be used to iterate through **Iterables** or you can call the **hasNext()** and **next()** functions.
- Usable data types: **int**, **float**, **String**. You may also use lists and custom data types composed of the aforementioned types.
- The method **intersection(list1, list2)** returns a list that is the intersection of list1 and list2.

1. Given the student's name and the course taken, output each student's name and total GPA.

**Declare any custom data types here:**

CourseData:

```
int courseID
float studentGrade // a number from 0-4
```

**map(String student, CourseData value):**  
emit( student, student grade)

**reduce( String key,**  
**Iterable< float > values):**

```
float sum = 0; int n = 0; float GPA;
```

2. Given a person's unique int ID and a list of the IDs of their friends, compute the list of mutual friends between each pair of friends in a social network.

**Declare any custom data types here:**

FriendPair:

```
int friendOne
int friendTwo
```

**map(int personID, list<int> friendIDs):**

```
for friendID in friendIDs: emit(pe
```

```
for friend in friendIDs: if personID < friendID:
```

**reduce( FriendPair key,**  
**Iterable< list<int> > values):**

```
mutualfriends = intersection(friendlists[1], friendlists[2]);
emit(FriendPair key, mutualfriends
```

```
mutualFriends = intersection( values.next(), values.next())
```

3. a) Given a set of coins and each coin's owner, compute the number of coins of each denomination that a person has.

**Declare any custom data types here:**

CoinPair:

String person

String coinType

**map(String person, String coinType):**

CoinPair = (person, coinType)

emit(CoinPair, 1)

**reduce(CoinPair key,  
Iterable< int > values):**

int sum = 0; for i in

b) Using the output of the first MapReduce, compute the amount of money each person has. The function `valueOfCoin(String coinType)` returns a float corresponding to the dollar value of the coin.

**map(CoinPair key, int amount):**

float value = valueOfCoin(key.coinType)

float sum = value\*amount;

emit(key.person, sum)

**reduce(String key,  
Iterable< float > values):**

int sum = 0; for i in

## Spark

- RDD: primary abstraction of a distributed collection of items
- Transforms: RDD  $\rightarrow$  RDD

<b>map</b> ( <i>func</i> )	Return a new distributed dataset formed by passing each element of the source through a function <i>func</i> .
<b>flatMap</b> ( <i>func</i> )	Similar to map, but each input item can be mapped to 0 or more output items (so <i>func</i> should return a Seq rather than a single item).
<b>reduceByKey</b> ( <i>func</i> )	When called on a dataset of (K,V) pairs, returns a dataset of (K,V) pairs where the values for each key are aggregated using the given reduce function <i>func</i> , which must be of type (V,V) $\Rightarrow$ V.

- Actions: RDD  $\rightarrow$  Value

<b>reduce</b> ( <i>func</i> )	Aggregate the elements of the dataset <i>regardless of keys</i> using a function <i>func</i>
-------------------------------	--

### 1. Implement Problem 1 of MapReduce with Spark

```
# students: list((studentName, courseData))
studentsData = sc.parallelize(students)
out = studentsData.map(lambda (k, v): (k, (v.studentGrade, ____)))
```

### 2. Implement Problem 2 of MapReduce with Spark

```
def genFriendPairAndValue(pID, fIDs):
    return [(pID, fID), fIDs] if pID < fID else (fID, pID) for fID in fIDs]
def intersection(l1, l2):
    return [x for x in l1 if x in l2]
# persons: list((personID, list(friendID))
personsData = sc.parallelize(persons)
```

### 3. Implement Problem 3 of MapReduce with Spark

```
# coinPairs: list((person, coinType))
coinData = sc.parallelize(coinPairs)
```